

Parallel and distributed computing

N. Kälin

March 7, 2021

Contents

1 Architectures	3
1.1 Implicit vs. explicit parallelism	3
1.2 Parallel programming models	3
1.3 Different grains of parallelism	4
1.4 Control structure of parallel platforms	5
1.5 Communication models of parallel platforms	6
1.6 Interconnection networks	6
1.7 Network topologies	8
1.8 Communication costs in parallel systems	13
2 Shared Memory Systems	15
2.1 Shared-address-space platforms	15
2.2 Cache coherence	16
2.3 Parallel programming in modern C++	17
2.4 Threads, Futures, Tasks	18
2.5 OpenMP	22
2.6 OpenMP: behind the scene	23
2.7 Concurrent tasks in OpenMP	25
2.8 Synchronization in OpenMP	28
2.9 OpenMP library functions	29

1 Architectures

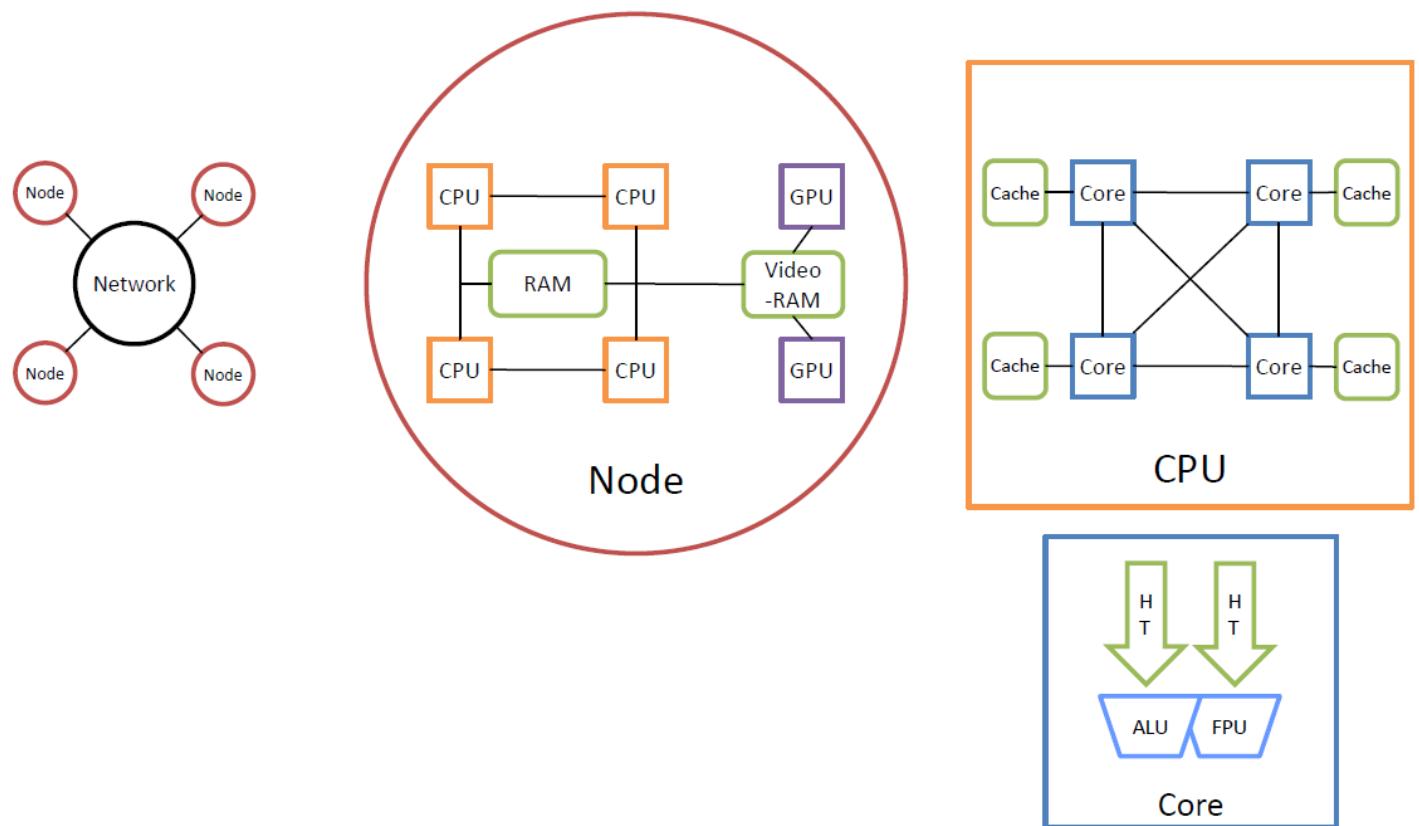


Figure 1: Parallel machine model (cluster)

1.1 Implicit vs. explicit parallelism

- Implicit Parallelism
 - processors have multiple functional units and execute multiple instructions in the same cycle
 - the precise way these instructions are selected and executed provides impressive diversity in Architectures
 - * **pipelining**
 - * **superscalar execution**
 - * **very long instruction word processors**
- Explicit Parallelism
 - an explicitly parallel program must specify concurrency (**control structure**) and interaction (**communication model**) between concurrent subtasks

1.2 Parallel programming models

1.2.1 Overview of Programming models

- Programming models
 - provide support for expressing concurrency and synchronization
- Process based models
 - assume that all data associated with a process is private, by default, unless otherwise specified
- Lightweight processes and Threads

- assume that all memory is global (bounded by process boundaries)
- memory protection between threads of the same process is not necessary
- support much faster memory access than processes with explicitly allocated shared memory
- Parallel programming language with syntax to specify parallelism
 - Examples: Ada, SR, Occam (no longer common)
- Directive based programming models
 - extend the threaded model by facilitating creation and synchronization of threads
 - Examples: Open MP, Linda, POP-C++

1.2.2 Parallel Machine Model

- PRAM
 - a natural extension of the Random Access Machine (RAM) serial architecture
 - consists of p processors and a global memory of unbounded size that is uniformly accessible to all processors
 - processors share a common clock but may execute different instructions in each cycle
- Handling of simultaneous memory accesses
 - Exclusive-read, exclusive-write (EREW)
 - Concurrent-read, exclusive-write (CREW)
 - Exclusive-read, concurrent-write (ERCW)
 - Concurrent-read, concurrent-write (CRCW)

What does concurrent write mean?

Common: write only if all values are identical.

Arbitrary: write the data from a randomly selected processor.

Priority: follow a predetermined priority order.

Sum: write the sum of all data items.

1.3 Different grains of parallelism

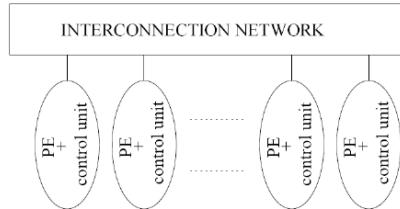
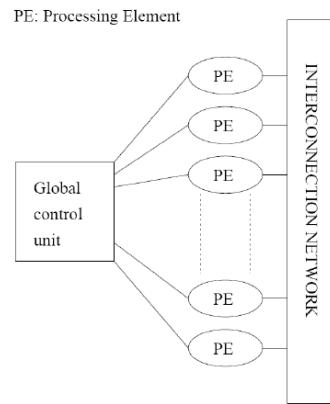
- Granularity: the ratio of computation to communication
 - periods of computation are separated from periods of communication by synchronization events
 - constrained by the inherent characteristics of the used algorithms
 - the parallel programmer must select the right granularity to benefit from the underlying platform
- Chunking
 - determining the amount of data to assign to each task (chunk or grain size)
- Which Granularity will lead to best performance?
 - depends on the algorithm and the used hardware environment
 - general rule: increase grain size if the communication overhead is too large

1.3.1 Trade-offs associated with chunk size

- Fine-grained parallelism
 - low arithmetic intensity
 - may not have enough work to hide long-duration asynchronous communication
 - facilitates load balancing by providing a larger number of more manageable (i.e. smaller) work units
 - too fine granularity can produce slower parallel implementation than the serial execution (too much overhead required for communication)
- Coarse-grained parallelism
 - high arithmetic intensity
 - complete applications can serve as the grain of parallelism
 - more difficult to load balance efficiently

1.4 Control structure of parallel platforms

- SIMD: Single Instruction stream, Multiple Data stream
 - there is a single control unit that dispatches the same instruction to various processors (that work on different data)
 - data parallelization
- MIMD: Multiple Instruction stream, Multiple Data stream
 - each processor has its own control unit
 - each processor can execute different instructions on different data items



1.4.1 SIMD Computers

- Hardware requirements
 - SIMD computers require less HW than MIMD computers (only one control unit)
 - SIMD computers require less memory (only one copy of the program is stored)
- Current implementations
 - Graphics Processing Units (GPUs)
 - Digital Signal Processors (DSPs) are widely used in cameras and sound equipments
 - Co-processing units in Intel CPUs: SSE_x, AVX-512
- Software requirements
 - SIMD relies on the regular structure of computations (such as those in image and video processing or in deep learning)
 - it is often necessary to selectively turn off operations on certain data items

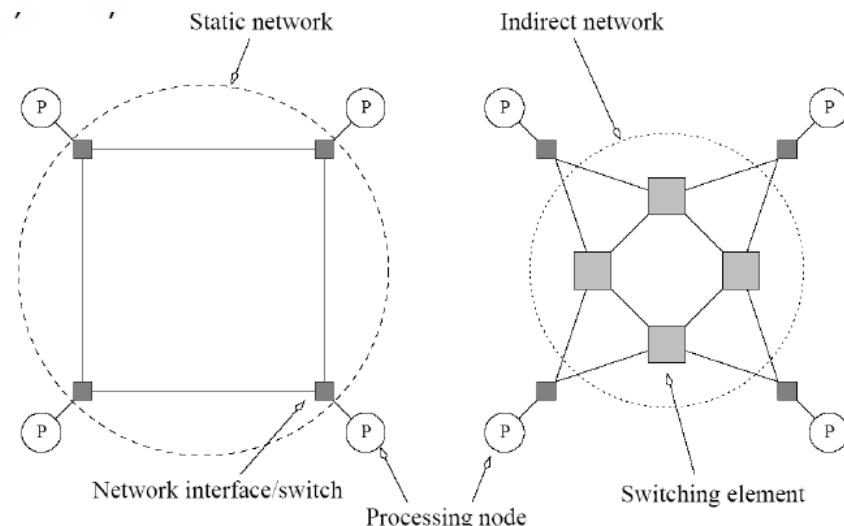
1.4.2 MIMD Computers

- Single Program Multiple Data (SPMD)
 - a simple variant of MIMD executes the same program on different processors
 - SPMD and MIMD are closely related in terms of programming flexibility and underlying architectural support
 - a single program consisting of several programs in a large switch block with conditions specified by the task identifiers is equivalent to the MIMD model
- Current MIMD implementations
 - SPARC servers, multiprocessor PCs, NASA Beowulf inspired workstation clusters
- Key advantages of workstation clusters
 - high performance workstations and PCs available at low cost
 - latest processors can easily be incorporated into the system as they become available
 - existing software can be used or modified

1.5 Communication models of parallel platforms

- Shared-Address-Space Platforms (Multiprocessors)
 - part (or all) of the memory is accessible to all processors
 - processors interact by modifying data objects stored in this shared-address-space
 - uniform or non-uniform memory access time (UMA vs. NUMA)
- Message Passing Platforms (Multicomputers)
 - comprise of a set of processors and their own (exclusive) memory
 - instances come naturally from clustered workstations (distributed systems) and non-shared-address-space multi-computers
 - are programmed using sending messages (variants of send and receive primitives)
 - libraries such as MPI (1990's) provide such primitives

1.6 Interconnection networks



- Interconnection Networks for Parallel Computers
 - carry data between processors and to memory
 - are made of switches and links (wires, fiber)
 - are classified as static or dynamic
 - * static (direct) networks consist of point-to-point communication links among processing nodes
 - * dynamic (indirect) networks are built using switches and communication links
- Network Topologies
 - a variety of network topologies have been proposed and implemented
 - tradeoff performance for cost
 - two basic categories: physical and logical topologies
 - commercial machines often implement hybrids of multiple topologies

1.6.1 Interconnection Network for HPC

- Infiniband
 - a computer-networking communications standard used in HPC that features very high throughput and very low latency
 - it is used for data interconnect both among and within computers
 - it is also utilized as either a direct, or switched interconnect between servers and storage systems, as well as an interconnect between storage systems
 - it is designed to be scalable and uses a switched fabric network topology

Year	FDR 2011	EDR 2014	HDR 2017	NDR 2020	XDR 2023
Throughput, per 1x [Gbit/s]	13.64	25	50	100	250
Speed for 4x links [Gbit/s]	54.54	100	200	400	1000
Speed for 12x links [Gbit/s]	163.64	300	600	1200	3000
Latency [μ s]	0.7	0.5	0.5	tbd	tbd

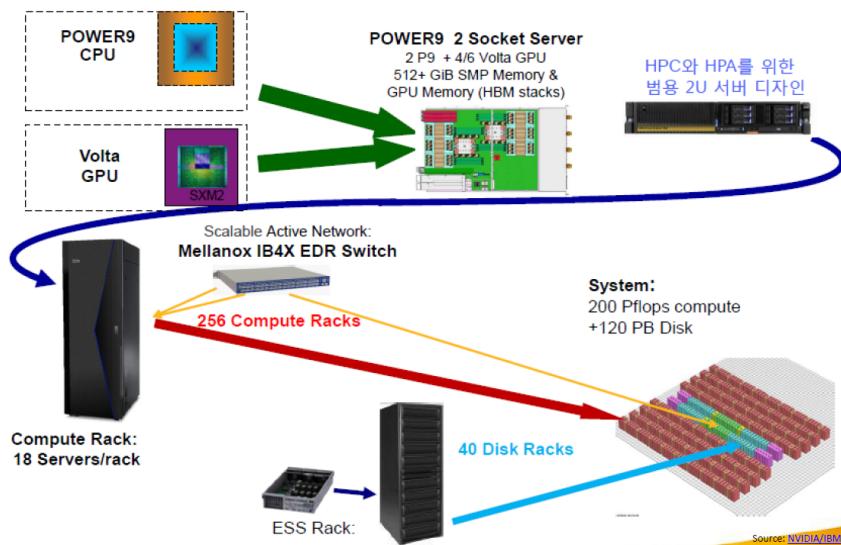
- PCI Express Version 4

- a high-speed serial computer expansion bus standard
- has numerous improvements over the older standards
 - * higher maximum system bus throughput
 - * lower I/O pin count and smaller physical footprint
- has been drafted with final specifications expected in 2017
- throughput:
 - * x1: 1.969 GByte/s
 - * x16: 31.508 GByte/s
- external cabling: Thunderbolt

- NVIDIA NVLink

- is a high-bandwidth, energy-efficient interconnect
- enables ultra-fast communication between the CPU and GPU, and between GPU
- throughput:
 - * version 1 (used in NVIDIA Pascal): x1: 20 GByte/s, x4: 80 GByte/s
 - * version 2 (used in IBM Power9 chip, NVIDIA Volta GPUs): x1: 25 GByte/s, x8: 200 GByte/s

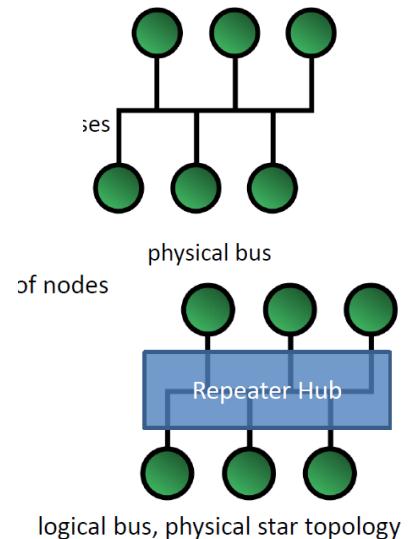
1.6.2 Data-Centric IT Environments



1.7 Network topologies

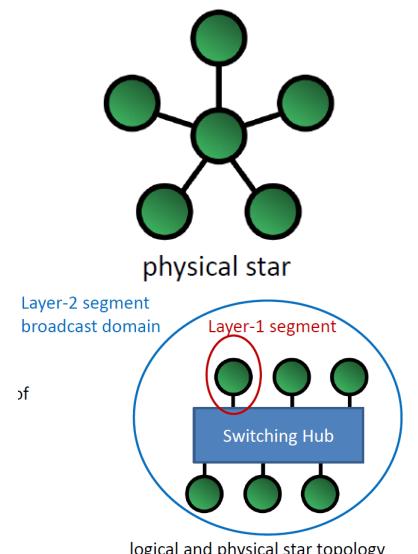
1.7.1 Network Topologies: Bus

- Principle and Properties
 - some of the simplest and earliest parallel machines used buses
 - all processors access a common bus for exchanging data
 - the distance between any two nodes is $O(1)$ in a bus
- Bottleneck
 - the bandwidth of the shared bus is a major bottleneck
 - typical bus-based machines are limited to dozens of nodes
- Examples
 - WLAN zone (logical bus topology)
 - PCI bus (physical bus topology)



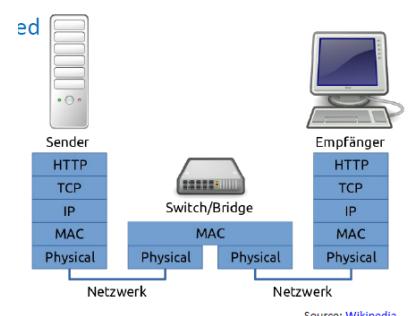
1.7.2 Network Topologies: Star

- Principle and Properties
 - every node is connected only to a common node at the center
 - distance between any pair of nodes is $O(1)$
- Bottleneck
 - the central node
- Example
 - today's Ethernet based LANs with bridging hub (Bridge) or switching hub (Switch) as the center of the star topology



Network Infrastructure: Switching Hub

- Principle and Properties
 - frame forwarding depends on learned physical device-addresses (MAC) per port (Layer-2 switching)
 - non-blocking: several input-output connections can be used in parallel without blocking
 - store-and-forward
 - * the switch buffers and verifies each frame before forwarding it
 - * a frame is received in its entirety before it is forwarded
 - * error checking can be done before forwarding
 - cut-through
 - * the switch starts forwarding after the frame's destination address is received
 - * when the outgoing port is busy at the time, the switch falls back to store-and-forward operation
 - * there is no error checking with this method

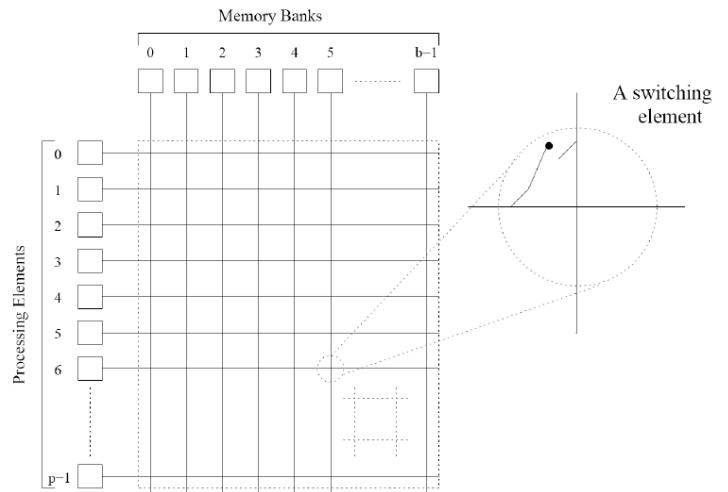


Switching Hub: Advanced Features

- Spanning Tree Protocol → Shortest Path Bridging
 - classic bridges may also interconnect using a spanning tree protocol that disables links so that the resulting local area network is a tree without loops
 - in contrast to routers, spanning tree bridges must have topologies with only one active path between two points
 - IEEE 802.1aq allows all paths to be active with multiple equal cost paths
 - * provides much larger layer 2 topologies (up to 16 million compared to the 4096 VLANs limit)
 - * improves the use of the **mesh topologies** through increased bandwidth and redundancy between all devices by allowing traffic to load share across all paths of a mesh network
- IEEE 802
 - is a family of IEEE standards dealing with local area networks and metropolitan area networks
 - services and protocols specified in IEEE 802 map to the lower two layers (Data Link and Physical) of the seven-layer OSI networking reference model
 - small subset of the working groups
 - * 802.1: higher layer LAN protocols (bridging)
 - * 802.1D: Spanning Tree Protocol (forwarding stopped while the spanning tree re-converged)
 - * 802.1S: Multiple Spanning Tree Protocol
 - * 802.1W: Rapid Spanning Tree Protocol
 - * 802.1aq: Shortest Path Bridging (SPB) (incorporate all the older spanning tree protocols)

1.7.3 Network Topologies: Crossbar

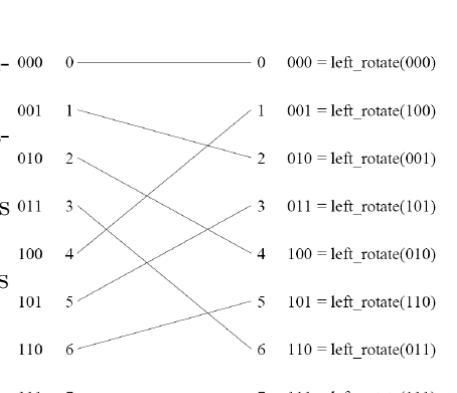
- Principle and Properties
 - a crossbar network uses an $p \cdot b$ grid of switches to connect p inputs to b outputs in a non-blocking manner
- Bottleneck
 - the cost of a crossbar of p processors grows as $O(p^2)$ → difficult to scale for large values of p
- Usage
 - in non-blocking switches
 - between L2- and L2-caches



1.7.4 Network Topologies: Multistage Network

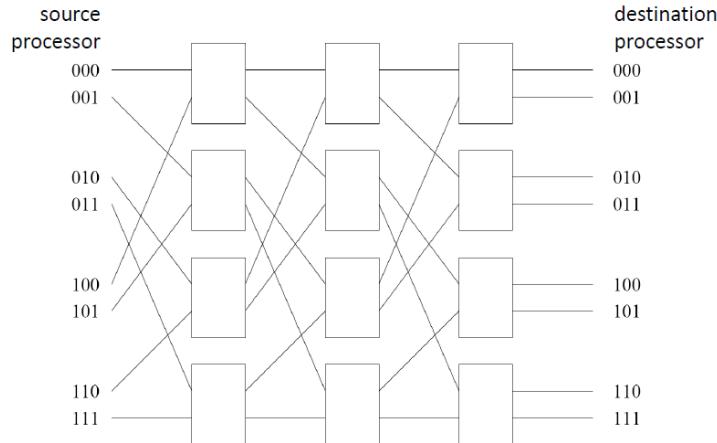
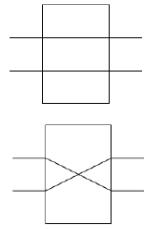
- Scalability
 - busses have excellent cost scalability, but poor performance scalability
 - crossbars have excellent performance scalability but poor cost scalability
 - multistage interconnects strike a compromise between these extremes
- Example: Omega Network
 - it consists of $\log(p)$ stages, where p is the number of inputs/outputs
 - at each stage, input i is connected to output j if:

$$j = \begin{cases} 2i, & 0 \leq i \leq p/2 - 1 \\ 2i + 1 - p, & p/2 \leq i \leq p - 1 \end{cases}$$



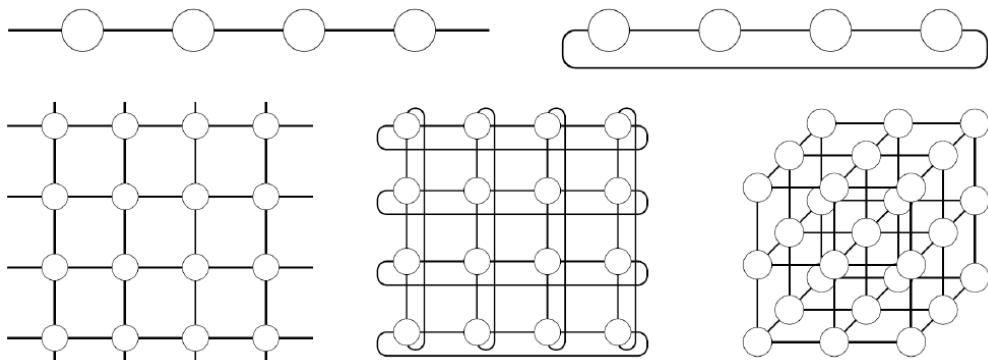
Omega Network

- Principle and Properties
 - the perfect shuffle patterns are connected using 2x2 switches
 - the switches operate in two modes: pass-through or cross-over



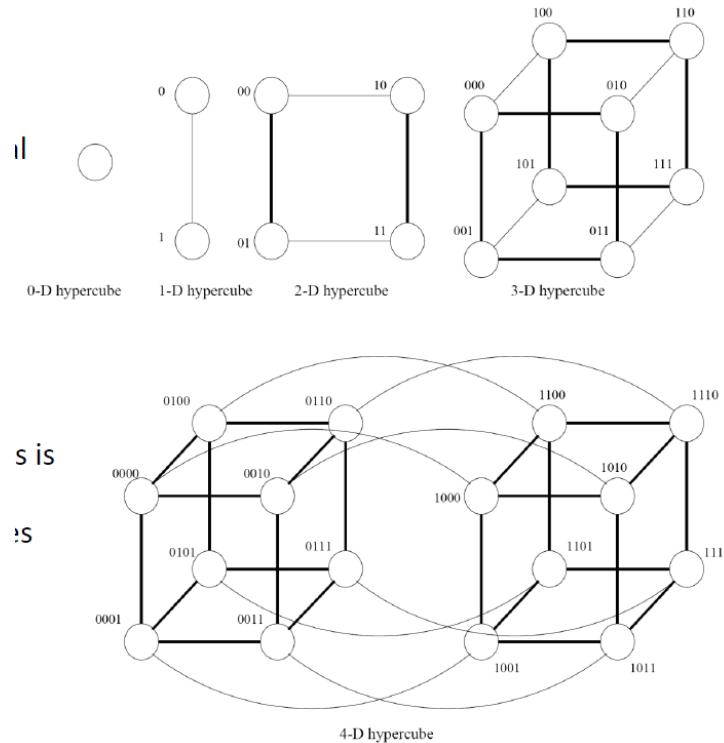
Linear Array, Mesh, and $k-d$ Mesh

- Principle and Properties
 - in a linear array, each node has two neighbors, one to its left and one to its right
 - if the nodes at either end are connected, we refer to it as a 1-D torus or a ring
 - a generalization to 2 dimensions has nodes with 4 neighbors, to the north, south, east, and west (toroidal mesh)
 - a further generalization to d dimensions has nodes with $2d$ neighbors

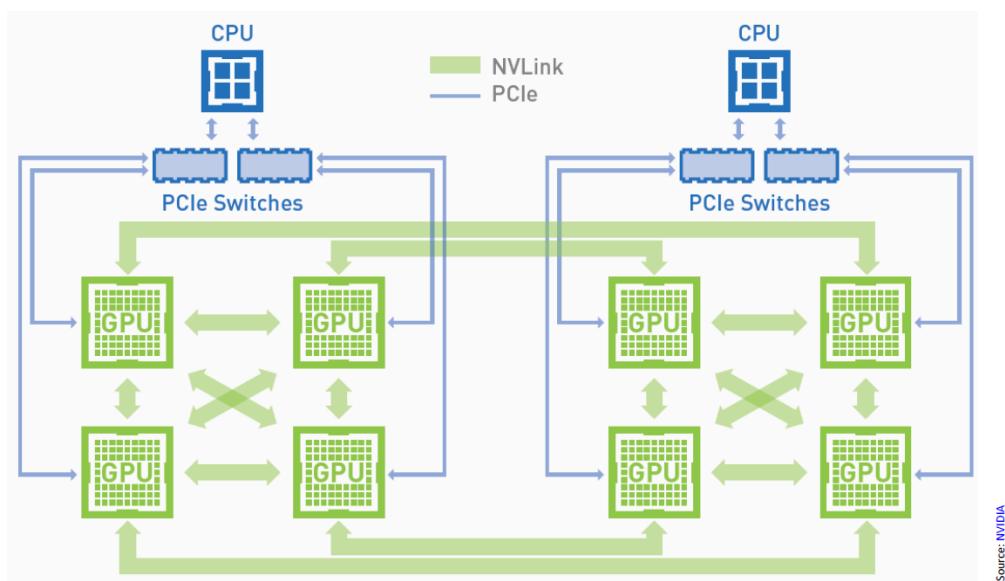


1.7.5 Network Topologies: Hypercube

- Principle and Properties
 - a special case of a d -dimensional mesh is a hypercube
 - $d = \log(p)$, where p is the total number of nodes
 - the distance between any two nodes is at most $\log(p)$
 - each node has $\log(p)$ neighbors
 - the distance between two nodes is given by the number of bit positions at which the two nodes differ



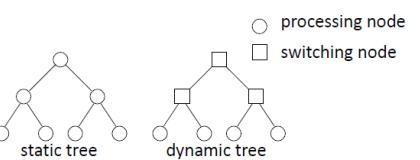
NVIDIA NVLink: Hypercube Mesh Hybrid



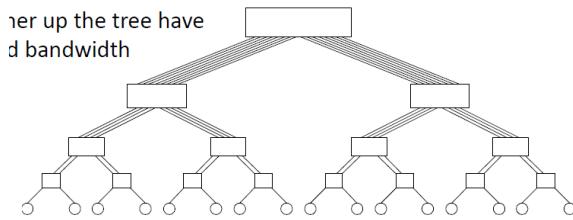
Source: NVIDIA

1.7.6 Tree-Based Networks

- Principle and Properties
 - one path between any pair of nodes
 - * linear arrays and star-connected networks are special cases of tree networks
 - the distance between any two nodes is no more than $2\log(p)$
 - links higher up the tree potentially carry more traffic than those at the lower levels
 - trees can be laid out in 2D with no wire crossings
- Fat-Tree



- links higher-up the tree have increased bandwidth



1.7.7 Evaluating Interconnection Networks

- Diameter
 - the distance between the farthest two nodes in the network
- Channel Bandwidth = channel width x channel rate
 - channel width: number of bits that can be communicated simultaneously over a link
 - channel rate: peak data transfer rate per link
- Cross-Section Bandwidth = bisection width x channel bandwidth
 - bisection width: the minimum number of wires one must cut to divide the network into two equal parts
- Cost
 - the number of links or switches (whichever is asymptotically higher) is a meaningful measure of the cost
 - the ability to layout the network
 - the length of wires
- ...

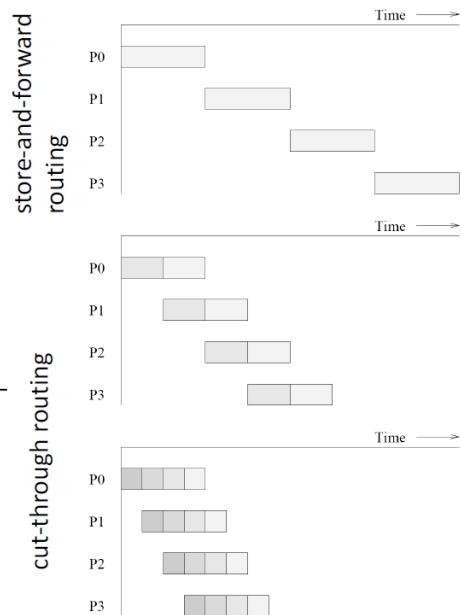
Network	Diameter	Bisection width	Arc connectivity	Cost (No. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$
Star	2	*	1	$p - 1$
Complete binary tree	$2 \cdot \log((p + 1)/2)$	1	1	$p - 1$
Linear array	$p - 1$	1	1	$p - 1$
2D Mesh, no wraparound	$2(\sqrt{p} - 1)$	\sqrt{p}	2	$2(p - \sqrt{p})$
2D wraparound Mesh	$2\lfloor\sqrt{p}/2\rfloor$	$2\sqrt{p}$	4	$2p$
Hypercube	$\log(p)$	$p/2$	$\log(p)$	$(p \cdot \log(p))/2$
Wraparound k -ary d -cube	$d\lfloor k/2 \rfloor$	$2k^{d-1}$	$2d$	dp

* depends on the node (switch) in the center, e.g. Crossbar or Omega Network

Network	Diameter	Bisection width	Arc connectivity	Cost (No. of links)
Crossbar	1	p	1	p^2
Omega Network	$\log(p)$	$p/2$	2	$p/2$
Dynamic Tree	$2 \cdot \log(p)$	1	2	$p - 1$

1.8 Communication costs in parallel systems

- Overhead in parallel programs
 - idling
 - contention
 - communication
- Communication costs depend on
 - communication model
 - the network topology
 - data handling and routing (e.g. packet routing, cut-through routing)
 - associated software protocols
 - ...



1.8.1 Message Passing Costs

- Total time to transfer a message over a network comprises of the following:
 - *Startup time (t_s)*: Time spent at sending and receiving nodes (executing the routing algorithm, programming routers, etc.)
 - *Per-hop time (t_h)*: This time is a function of number of hops and includes factors such as switch latencies, network delays, etc.
 - *Per-word transfer time (t_w)*: This time includes all overheads that are determined by the length of the message. This includes bandwidth of links, error checking and correction, etc.

1.8.2 Cost Model for Communicating Messages

- Communication Costs
 - the cost of communicating a message between two nodes/hops away using cut-through routing is given by

$$t_{\text{comm}} = t_s + l \cdot t_h + m \cdot t_w$$
 - t_h is typically smaller than t_s and t_w , so the second term does not show, when m is large
 - furthermore, it is often not possible to control routing and placement of tasks
- Simplified Cost Model

$$t_{\text{comm}} = t_s + m \cdot t_w$$
- Remarks
 - it is important to note that the original expression for communication time is valid for only uncongested networks
 - if a link takes multiple messages, the corresponding t_w term must be scaled up by the number of messages
 - different communication patterns congest different networks to varying extents

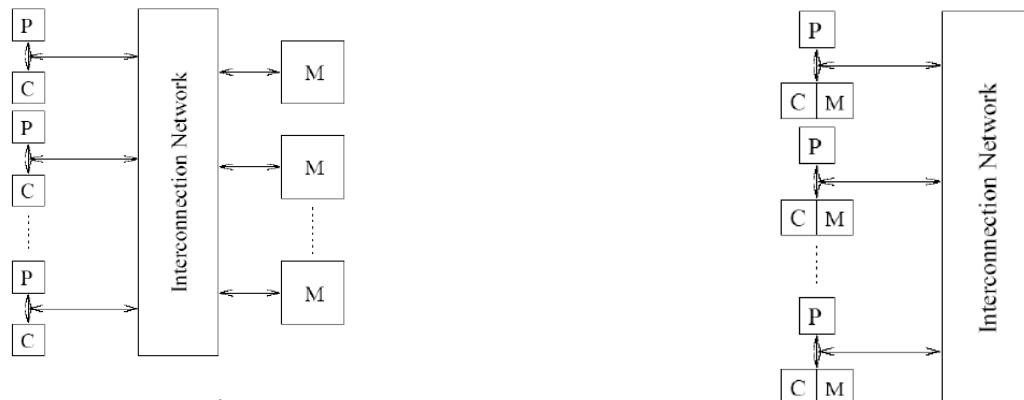
1.8.3 Cost Model for Shared Memory Systems

- Simplified Cost Model (still practical, but accurate cost modeling is more difficult)
 - memory layout is typically determined by the system
 - finite cache sizes can result in cache thrashing

- overheads associated with invalidate and update operations are difficult to quantify
- spatial locality is difficult to model
- pre-fetching can play a role in reducing the overhead associated with data access
- false sharing and contention are difficult to model

2 Shared Memory Systems

2.1 Shared-address-space platforms



UMA (Uniform-memory-access)

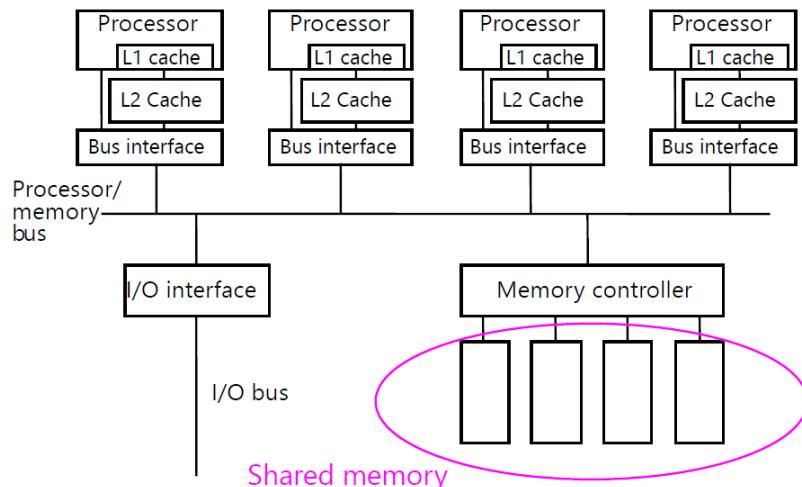
shared-address-space computer with local caches and
global memories → all memory access times (except cache) are identical

NUMA (Non-uniform-memory-access)

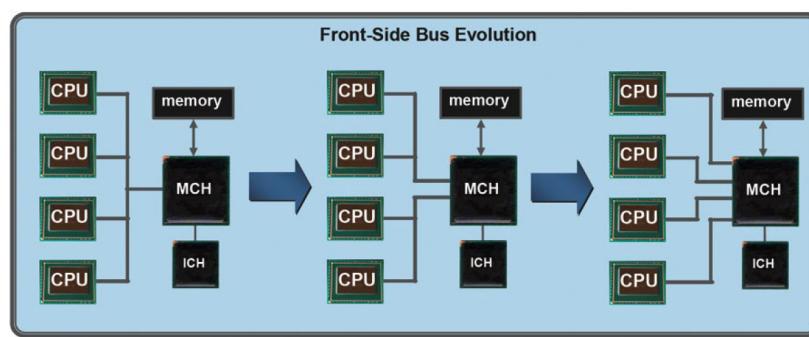
shared-address-space computer with local memory only
→ local memory access times are shorter

2.1.1 UMA Examples

- Intel Front Side Bus Architecture



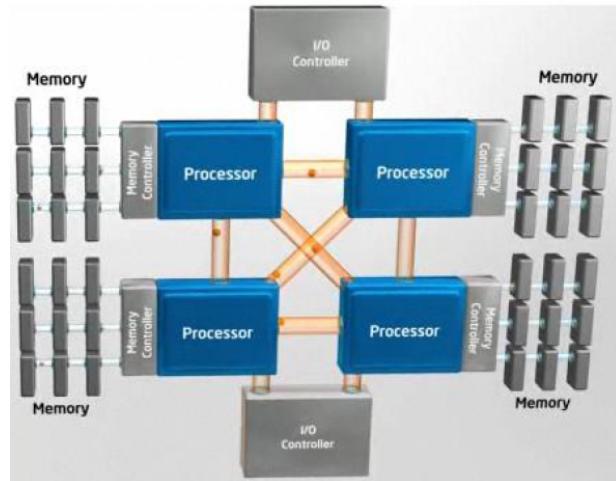
- Intel Pentium Front Side Bus Evolution



Source: Hardware LUXX

2.1.2 NUMA Examples

- Example: Intel Core i7



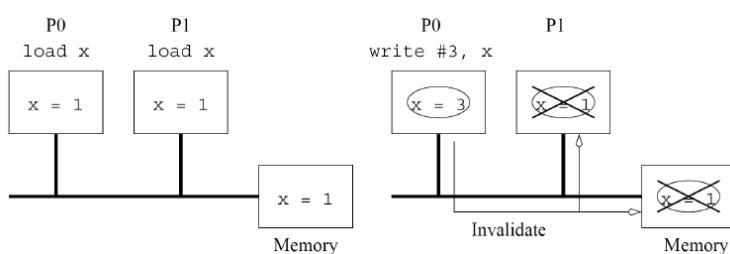
[Source: University of Portsmouth](#)

2.2 Cache coherence

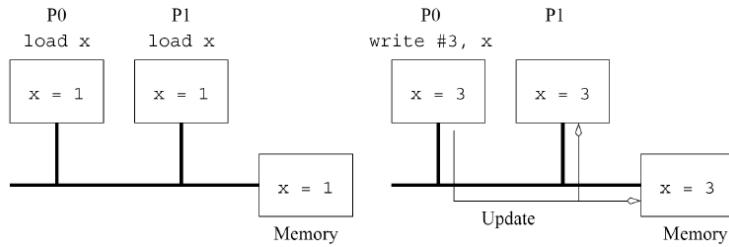
- Cache
 - principle of space and time locality
 - faster memory access
 - additional hardware is required to keep multiple copies of data consistent with each other
- Cache Coherence
 - ensuring that concurrent operations on multiple copies of the same memory word have well-defined semantics
 - this semantic is generally one of serializability
 - * there exists some serial order of instruction execution that corresponds to the parallel schedule

2.2.1 Update and Invalidate Protocols

- Scenario
 - when a processor changes the value of its copy of a variable, one of two things must happen:
 - * the other copies must be invalidated (invalidate protocol) (write-back)



* the other copies must be updated (update protocol) (write-through)



- Pros and Cons

- if a processor just reads a value once and does not need it again, an update protocol may generate significant overhead
- if two processors make interleaved test and updates to a variable, an update protocol is better
- both protocols suffer from false sharing overheads (two words that are not shared, however, they lie on the same cache line)

- Most current machines use invalidate protocols

- each copy of a data item is associated with a state (e.g. shared, invalid, or dirty)
 - * in shared state, there are multiple valid copies of the data item (therefore, an invalidate would have to be generated on an update)
 - * in dirty state, only one copy exists and therefore, no invalidates need to be generated
 - * in invalid state, the data copy is invalid, therefore, a read generates a data request (and associated state changes)

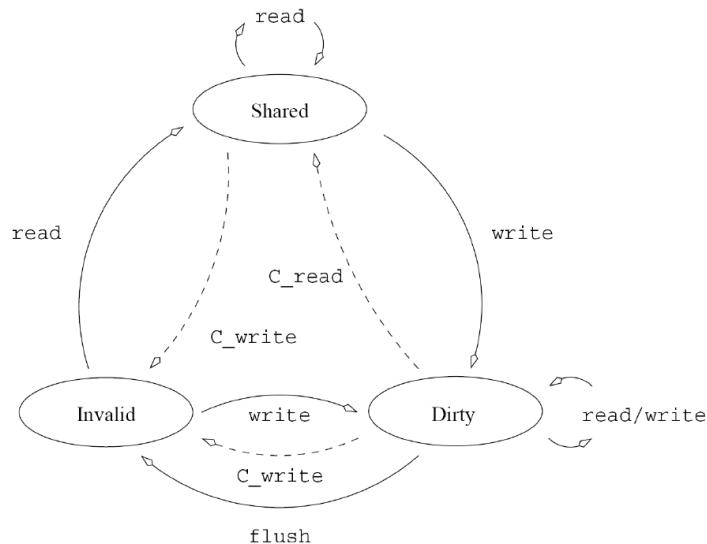


Figure 2: State Diagram of an Invalidate Protocol

2.3 Parallel programming in modern C++

- Explicit Threading
 - data exchange is more apparent
 - * this helps in alleviating some of the overheads from data movement, false sharing, and contention
 - provides richer API in the form of condition waits, locks of different types, and increased flexibility for building composite synchronization operations
 - tools and support are easier to find
- Directives Layered on Top of Threads
 - simplify a variety of thread-related tasks
 - a programmer is rid of the tasks of initializing attributes objects, setting up arguments to threads, partitioning iteration spaces, etc.

2.3.1 Thread

- A thread is a single stream of control in the flow of a program. A program like

```
for (row = 0; row < n; row++)
    for (column = 0; column < n; column++)
        c [row] [column] = dot_product (get_row (a, row), get_col (b, col));
```

- can be transformed to

```
for (row = 0; row < n; row++)
    for (column = 0; column < n; column++)
        thread ([&]) {
            c [row] [column] = dot_product (get_row (a, row), get_col (b, col));
        };
```

- In this case, one may think of the thread as an instance of a function that returns before the function has finished executing.

2.4 Threads, Futures, Tasks

2.4.1 C++: Threads and futures

- Thread
 - low-level
 - exchange of data must be synchronized itself
 - uncaught exceptions in the thread function lead to the termination of the entire program
 - `thread_local` storage class: static/global variables are created for each thread
 - is started automatically in the constructor
- Future
 - asynchronous processing: parallel execution or when calling `get()`
 - exceptions appear in the parent thread when the result is picked up with `get()`
 - when leaving the scope of the responsible future, then its destructor ensures a clear and problem-free termination of the asynchronous computation

2.4.2 C++: Threads as the Basis of Parallelism

- Constructor and Executor
 - `thread(executable object, parameters of the executable object)`
 - Executable object
 - * function object (functor)
 - * lambda expression
 - * pointer to a function
 - the executable object and the parameters are copied by default, so the thread can work on their own data
- Example

```

void printFibs(size_t from,
    size_t to);
{
    struct Image {
        void fill(int r, int g, int b);
    };
}

int main() {
    thread th1(printFibs, 28, 35);
    Image img;
    thread th2([&img] { img.fill(0, 1, 2); });
    th1.join(); th2.join();
}

```

2.4.3 C++: Futures and async

- Asynchronous Computation
 - parallel or deferred processing
 - `async()` initiates a computation and returns immediately
 - `get()` blocks until the result of the computation is available
- Return value of `async()`
 - `future<RT>`, where RT is the return type of the asynchronously executed function
- Launch Policy
 - `async(launch::async, longComputation)` guarantees parallel execution (default)
 - `async(launch::deferred, computation)` executed when calling `get()`
- Behind the scenes
 - a future can be produced without calling `async()`, by first creating a promise (some kind of transmission channel)

C++: Futures Example

```

#include <future>

static size_t fibrec(size_t n) {
    return (n < 2) ? 1 : fibrec(n - 2) + fibrec(n - 1);
}

int main() {
    // asynchronous computation
    auto fut1 = async(launch::async, &fibrec, 35);
    // deferred computation
    auto fut2 = async(launch::deferred, &fibrec, 35);
    cout << fut2.get() << endl;      // waiting for the result of fut2
    cout << fut1.get() << endl;      // waiting for the result of fut1
}

```

What happens if the order of the two calls of `get()` is changed? → `fut2` would only start after `fut1` would have finished → serial execution

2.4.4 C++: Packaged Tasks

- Concept
 - a `packaged_task` wraps a callable element and allows its result to be retrieved asynchronously
 - it is similar to function, but transferring its result automatically to a future object
- Syntax
 - `template<class Ret, class ... Args> class packaged_task<Ret(Args...)>;`
- Object contains internally two elements

- a stored task, which is some callable object whose call signature shall take arguments of the types in Args... and return a value of type Ret
- a shared state, which is able to store the results of calling the stored task (of type Ret) and be accessed asynchronously through a future

C++: Packaged Tasks Example

```
// create task for calling fibrec
// argument of fibrec has to be defined later
packaged_task<size_t(size_t)> task1(&fibrec);
auto f1 = task1.get_future(); // future for getting result

// create task for calling fibrec
// argument of fibrec is bound to 35
packaged_task<size_t(void)> task2(bind(&fibrec, 35));
auto f2 = task2.get_future(); // future for getting result

// call task1 in a parallel thread (move semantic)
thread th(move(task1), 35);
// call task2 in this thread
task2();

// get results
cout << f1.get() << endl;
cout << f2.get() << endl;

th.join(); // this thread waits on parallel thread th
```

2.4.5 C++: Synchronization Primitives

- Synchronized data access (read and write) is necessary if at least one of the parallel threads changes common data
- Synchronization Primitives
 - atomic_xyx all accesses are atomic (are not interrupted)
 - atomic_flag atomic bool but lock-free
 - once_flag used in call_once, makes sure that only one of the parallel threads executes the function
 - mutex realizes mutual exclusion
 - recursive_mutex allows a thread computing a recursive function to reenter a critical section
 - lock_guard locks a critical section; very simple usage; the only state is locked
 - unique_lock needs its unique mutex object; handles both states: locked and unlocked
 - condition_variable blocks this thread until signaled (signals or notifications can be sent by other threads)

2.4.6 Serial for- vs. Parallel for-Loop

- Notice
 - a lot of programming languages don't have a special parallel for-loop syntax
 - they use the keyword `for` in two situations
- Serial for-Loop
 - the loop notation simplifies the programming of a fixed number of repetitive, sequentially ordered steps
 - * Example: reading n integers from a sequential file
- Parallel for-Loop

- the loop notation simplifies the programming of a fixed number of repetitive steps, that can be done in any order
 - * Example: for each element of an array proceed the same task

C++: Possible parallel for-each implementation

```
template<typename It>
void parallelForEach(It begin, It end, function<void(typename It :: reference)> f) {
    const ptrdiff_t len = end - begin;
    if (len == 0) {
        return;
    } else if (len == 1) {
        f(*begin);
        return;
    }

    const It mid = begin + (ptrdiff_t)(len/2);
    future<void> fut = async(parallelForEach<It>, begin, mid, f);
    try {
        parallelForEach(mid, end, f);
    } catch (...) {
        fut.wait();
        throw;
    }
    fut.get();
}
```

C++: Usage of parallelForEach

```
static size_t fibrec(size_t n) {
    return (n < 2) ? 1 : fibrec(n - 2) + fibrec(n - 1);
}

int main() {
    int vals[] = {5, 10, 15, 20, 25, 30, 35, 40};
    int size = sizeof(vals)/sizeof(int);

    parallelForEach(vals, vals + size, [](size_t n) {
        cout << fibrec(n) << endl;
    });
    return 0;
}
```

2.4.7 Parallel Algorithms in C++17

- Ordering
 - "sequenced-before" is an asymmetric, transitive, pair-wise relationship between evaluations within the same thread
 - A is sequenced before B



- if A is not sequenced before B and B is not sequenced before A, then two possibilities exist:
 - * evaluations of A and B are **indeterminately sequenced**: they may be performed in any order but may not overlap

A B or B A
 - * evaluations of A and B are **unsequenced**: they may be performed in any order and may overlap (within a single thread of execution, the compiler may interleave the CPU instructions that comprise A and B)

A B or B A

- Execution Policies

- most algorithms have overloads that accept execution policies:
 - `std::execution::seq` sequential execution like calling the algorithms without an execution policy
 - `std::execution::par` execution potentially using multiple threads
 - parallel instructions are indeterminately sequenced
 - is not allowed to cause data races or to cause dead-locks
 - `std::execution::par_unseq` execution may be parallelized, vectorized, or migrated across threads
 - parallel instructions and ordering in the same thread: unsequenced
 - use of blocking synchronization primitives (e.g. mutex) may cause dead-lock
 - `std::execution::unseq` execution may be vectorized
 - ordering in the same thread: unsequenced

Execution Policies and Ordering Examples

```

int a[] = {0, 1};
vector<int> v;
for_each(execution::par, begin(a), end(a), [&](int i) {
    v.push_back(i*2+1); // Error: data race (vector isn't thread safe)
});

int x = 0;
mutex m;
for_each(execution::par, begin(a), end(a), [&](int) {
    lock_guard<mutex> guard(m);
    ++x; // Correct, because mutual exclusion is guaranteed
}); // and sequence among parallel lambdas is irrelevant

for_each(execution::par_unseq, begin(a), end(a), [&](int) {
    lock_guard<mutex> guard(m); // Error: calls m.lock() and several of these
    ++x; // calls are unsequenced and can interleave
})

```

2.5 OpenMP

- A Standard for Directive Based Parallel Programming
 - OpenMP is a directive-based API that can be used with
 - * FORTRAN
 - * C/C++
 - for programming shared address space machines
 - OpenMP directives provide support for
 - concurrency
 - synchronization
 - data handling

- while obviating the need for explicitly setting up
 - mutexes
 - condition variables
 - data scope
 - initialization
- standard specifications: <http://www.openmp.org/specifications>
 - Visual Studio 2019 only supports OpenMP standard 2.0

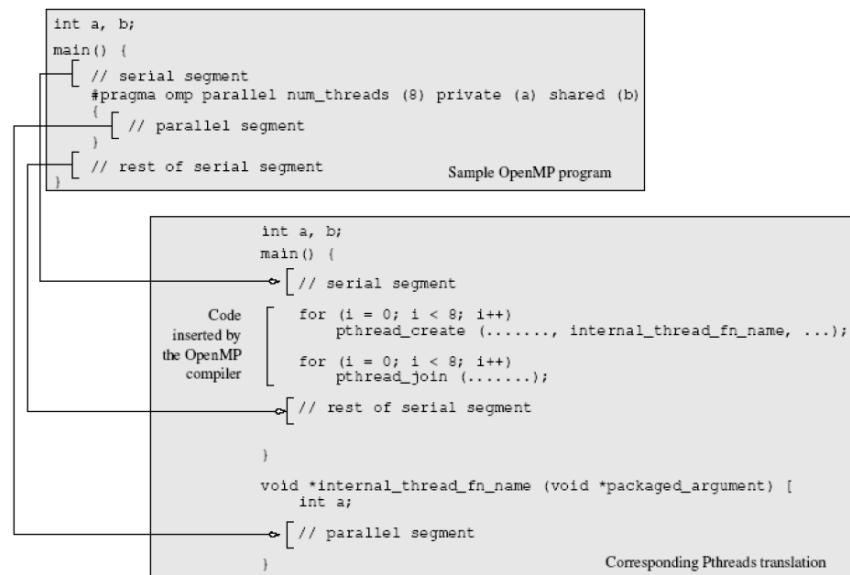
2.5.1 OpenMP Programming Model

- OpenMP in C/C++
 - directives are based on the `#pragma` compiler directives
 - a directive consists of a directive name followed by clauses
`#pragma omp directive [clause list]`
- OpenMP programs
 - execute serially until they encounter the parallel directive, which creates a group of threads

```
#pragma omp parallel [ clause list ]
/* structured block */
```

- the main thread that encounters the parallel directive becomes the **master** of this group of threads and is assigned the thread **id 0** within the group
- at the end of the parallel executed block the main thread waits for all parallel threads (join)

2.6 OpenMP: behind the scene



2.6.1 Clause List in OpenMP

- Clause List specifies
 - Conditional Parallelization
 - * the clause `if` (scalar expr) determines whether the parallel construct results in creation of threads
 - * the scalar expression is evaluated at runtime
 - Degree of Concurrency
 - * the clause `num_threads(integer expr)` specifies the number of threads that are created
 - Data Handling
 - * the clause `private` (variable list) indicates variables local to each thread T

- * the clause `firstprivate` (variable list) is similar to the `private`, except values of variables are initialized to corresponding values before the parallel directive
- * the clause `shared` (variable list) indicates that variables are shared across all the threads

- Example

```
#pragma omp parallel if(is_parallel == 1) num_threads(8) \
private(a) shared(b) firstprivate(c)
{
    /* structured block */
}
```

2.6.2 Default Clause in OpenMP

- Syntax
 - `default(shared — none)`
- Semantic
 - the default clause allows the user to affect the data-sharing attributes of variables
 - omitting this clause is the same as the `default(shared)`
- `default(shared)`
 - is equivalent to explicitly listing each currently visible variable in a `shared` clause, unless it is **thread-private** or **const**-qualified
- `default(None)`
 - it is usually better style to use `default(None)` instead of `default(shared)`
 - requires that at least one of the following must be true for every reference to a variable in the lexical extent of the parallel construct
 - * the variable is explicitly listed in a data-sharing attribute clause
 - * the variable is declared within the parallel construct
 - * the variable is **threadprivate** or has a **const**-qualified type
 - * the variable is the loop control variable for a **for**-loop that immediately follows a **for** or **parallel for**-directive, and the variable reference appears inside the loop

2.6.3 Default Clause in OpenMP

- Reduction Clause
 - specifies how multiple local copies of a variable at different threads are combined into a single copy at the master when threads exit
 - the usage is


```
reduction (operator: variable list)
```
 - the variables in the list are implicitly specified as being private to threads
 - the operator can be one of `+`, `*`, `-`, `&`, `|`, `^`, `&&`, `||`
- Example

```
#pragma omp parallel default(None) reduction(+: sum) num_threads(8)
{
    /* compute local sums here */
    sum = ...;
}
/* sum here contains sum of all local instances of sums */
```

2.6.4 OpenMP Programming: Example

```
int main() {
    const int npoints = 10000000;
```

```

int sum = 0;
mt19937_64 re; // random engine
uniform_real_distribution<double> dist;
#pragma omp parallel default(none) reduction(+: sum) num_threads(8)
{
    #pragma omp for
    for (int i = 0; i < npoints; i++) {
        if (hypot(dist(re), dist(re)) < 1) sum++
    }
}
cout << setprecision(10) << 4.0*sum/npoints << endl;
}

```

2.6.5 OpenMP Directives

Directive	Description
atomic	Specifies that a memory location will be updated atomically.
barrier	Synchronizes all threads in a team; all threads pause at the barrier, until all threads execute the barrier.
critical	Specifies that code is only executed on one thread at a time.
flush	Specifies that all threads have the same view of memory for all shared objects.
for	Causes the work done in a for-loop inside a parallel region to be divided among threads.
master	Specifies that only the master thread should execute a section of the program.
ordered	Specifies that code inside a parallelized for-loop should be executed by multiple threads in parallel.
sections	Identifies code sections to be divided among all threads.
single	Specifies that section of code should be executed on a single thread, not necessarily the master thread.
threadprivate	Specifies that a variable is private to a thread.

2.7 Concurrent tasks in OpenMP

- The `for`-directive
 - specifies concurrent iterations
 - is used to split parallel iteration spaces across threads
 - the general form is


```
#pragma omp for [ clause list ]
/* for loop */
```
 - allowed clauses
 - * `private`
 - * `firstprivate`, `lastprivate`
 - * `reduction`
 - * `schedule`
 - * `nowait`
 - * `ordered`

- The `sections`-directive
 - specifies concurrent tasks
 - the general form is


```
#pragma omp sections [ clause list ]
{
    #pragma omp section
        /* structured block 1 */
    #pragma omp section
        /* structured block 2 */
    // ...
}
```

2.7.1 The `for` Directive

- Schedule Clause

- deals with the assignment of iterations to threads
- the general form of the schedule directive is
schedule(scheduling_class [, parameter])
- four scheduling classes
 - static
 - * splits the iteration space into equal chunks of size *parameter* and assigns them to threads in a round-robin fashion
 - * when no *parameter* is specified, the iteration space is split into equally sized chunks, one chunk per thread
 - dynamic
 - * the iteration space is partitioned into chunks of size *parameter* (default value: 1)
 - * these chunks are assigned to threads as they become idle
 - guided
 - * the chunk size is reduced exponentially as each chunk is dispatched to a thread
 - * the *parameter* specifies the smallest chunk size (default value: 1)
 - runtime
 - * the environment variable OMP_SCHEDULE determines at runtime the scheduling class and the chunk size

2.7.2 The schedule Clause: Example

```
/* static scheduling of matrix multiplication loops */
#pragma omp parallel default(none) shared(a, b, c) num_thread(4)
#pragma omp for schedule(static)
for (int i = 0; i < 128; i++) {
    for (int j = 0; i < 128; j++) {
        c(i, j) = 0;
        for (int k = 0; k < 128; k++) {
            c(i, j) += a(i, k)*b(k, j);
        }
    }
}
```

2.7.3 The nowait Clause

- Implicit Barrier
 - at the end of the parallel `for`-loop all threads join
- Clause nowait
 - often, it is desirable to have a sequence of `for`-directives within a parallel construct that do not execute an implicit barrier at the end of each `for` directive
- Example

```
#pragma
{
    #pragma omp for nowait
    for (int i = 0; i < nmax; i++)
        if (isEqual(name, current_list[i])) processCurrName(name);
    #pragma omp for
    for (int i = 0; i < nmax; i++)
        if (isEqual(name, past_list[i])) processPastName(name);
}
```

2.7.4 The sections Directive

- Sections
 - supports non-iterative parallel task assignment using the sections directive
- Example

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            taskA();
        }
        #pragma omp section
        {
            taskB();
        }
        #pragma omp section
        {
            taskC();
        }
    }
}
```

2.7.5 Merging Directives

- Remember
 - the **parallel** directive creates the group of threads
 - the **for** and the sections directive would execute serially (by the master thread) if no parallel directive is specified before
- Merging parallel and **for** directives

```
#pragma omp parallel shared(n)
{
    #pragma omp for
    for (int i = 0; i < n; i++) {
        // ...
    }
}
```

```
#pragma omp parallel for shared(n)
for (int i = 0; i < n; i++) {
    // ...
}
```

2.7.6 Nesting parallel Directives

- Example

```
#pragma omp parallel for shared(a, b, c) num_threads(4)
for (int i = 0; i < 128; i++) {
    #pragma omp parallel for shared(a, b, c) num_threads(4)
    for (int j = 0; j < 128; j++) {
        c(i, j) = 0;
    }
    #pragma omp parallel for shared(a, b, c) num_threads(4)
    for (int k = 0; k < 128; k++) {
```

```

    c( i , j ) + = a( i , k ) * b( k , j ) ;
}
}
}
```

- Remarks

- OpenMP does not allow nested `for`, sections, and single directives that bind to the same parallel directive
 - each `for` directive brings its own parallel directive, which only generates a logical team of threads on encountering a nested parallel directive
 - the newly generated logical team is still executed by the same thread corresponding to the outer parallel directive
 - to generate a new set of threads, nested parallelism must be enabled by setting the `OMP_NESTED` environment variable to `TRUE`

2.8 Synchronization in OpenMP

- Synchronization Constructs

- `#pragma omp barrier`
 - `#pragma omp single [clause list]`
 `/* structured block */`
 - `#pragma omp master`
 `/* structured block */`
 - `#pragma omp critical [(name)]`
 `/* structured block */`
 - `#pragma omp atomic`
 `/* memory update instruction */`
 - `#pragma omp ordered`
 `/* structured block */`
 - `#pragma omp flush [(variable list)]`

2.8.1 Example: Prefix Sums

```
cumulSum[0] = list[0];
#pragma omp parallel for default(none) shared(cumulSum, list) ordered
for (int i = 1; i < n; i++) {
    // other work
#pragma omp ordered
{
    cumulSum[i] = cumulSum[i - 1] + list[i];
}
}
```

2.8.2 Data Handling in OpenMP

- Which data class should you use when?

- **private**
 - * a thread initializes and uses a variable and no other thread accesses the data
 - * it is better to use a local variable in an omp block
 - * if multiple threads manipulate different parts of a large data structure, the programmer should explore ways of breaking it into smaller data structures and making them **private** to the thread that manipulates them

- firstprivate
 - * a thread repeatedly reads a variable that has been initialized earlier in the program
- reduction
 - * if multiple threads manipulate a single piece of data, one must explore ways of breaking these manipulations into local operations followed by a single global operation (reduction)
- threadprivate(variable list)
 - * all variables in the list are local to each thread and are initialized once before they are accessed in a parallel region
 - * these variables persist across different parallel regions provided dynamic adjustment of the number of threads is disabled and the number of threads is the same

2.9 OpenMP library functions

- Thread and Processor Count
 - `#include <omp.h>`
 - `void omp_set_num_threads(int num_threads);`
 - `int omp_get_num_threads();`
 - `int omp_get_max_threads();`
 - `int omp_get_thread_num();`
 - `int omp_get_num_procs();`
 - `int omp_in_parallel();`
- Controlling and Monitoring Thread Creation
 - `void omp_set_dynamic(int dynamic);`
 - `int omp_get_dynamic();`
 - `void omp_set_nested(int nested);`
 - `int omp_get_nested();`
- Mutual Exclusion
 - `void omp_init_lock(omp_lock_t *lock);`
 - `void omp_destroy_lock(omp_lock_t *lock);`
 - `void omp_set_lock(omp_lock_t *lock);`
 - `void omp_unset_lock(omp_lock_t *lock);`
 - `int omp_test_lock(omp_lock_t *lock);`

2.9.1 Environment Variables in OpenMP

- Environment Variables
 - OMP_NUM_THREADS
 - * specifies the default number of threads created upon entering a parallel region
 - OMP_SET_DYNAMIC
 - * determines if the number of threads can be dynamically changed
 - OMP_NESTED
 - * turns on nested parallelism
 - OMP_SCHEDULE
 - * scheduling of `for`-loops if the clause specifies runtime
- Common Mistakes in OpenMP programs
 - http://michaelsuess.net/publications/suess_leopold_common_mistakes_06.pdf

2.9.2 OpenMP Standard

- Version 3.1
 - major change since version 2.5
 - * tasks added to execution model
- Version 4.0

- major changes since version 3.1
 - * array syntax extended to support array sections
 - * proc_bind clause to support thread affinity policies
 - * SIMD constructs to support SIMD parallelism
 - * device constructs to support execution of devices (e.g. GPU)
 - * user defined reductions
 - * depend clause to support task dependencies
- Version 5.1
 - major changes since version 4.5
 - * extended memory model to distinguish different types of flush operations
 - * support of modern C++20