

Parallel and distributed computing

N. Kälin

March 6, 2021

Contents

1 Architectures	3
1.1 Implicit vs. explicit parallelism	3
1.2 Parallel programming models	3
1.3 Different grains of parallelism	4
1.4 Control structure of parallel platforms	5
1.5 Communication models of parallel platforms	6
1.6 Interconnection networks	6
1.7 Network topologies	8
1.8 Communication costs in parallel systems	13

1 Architectures




Figure 1: Parallel machine model (cluster)

1.1 Implicit vs. explicit parallelism

- Implicit Parallelism
 - processors have multiple functional units and execute multiple instructions in the same cycle
 - the precise way these instructions are selected and executed provides impressive diversity in Architectures
 - * **pipelining**
 - * **superscalar execution**
 - * **very long instruction word processors**
- Explicit Parallelism
 - an explicitly parallel program must specify concurrency (**control structure**) and interaction (**communication model**) between concurrent subtasks

1.2 Parallel programming models

1.2.1 Overview of Programming models

- Programming models
 - provide support for expressing concurrency and synchronization
- Process based models
 - assume that all data associated with a process is private, by default, unless otherwise specified
- Lightweight processes and Threads

- assume that all memory is global (bounded by process boundaries)
- memory protection between threads of the same process is not necessary
- support much faster memory access than processes with explicitly allocated shared memory
- Parallel programming language with syntax to specify parallelism
 - Examples: Ada, SR, Occam (no longer common)
- Directive based programming models
 - extend the threaded model by facilitating creation and synchronization of threads
 - Examples: Open MP, Linda, POP-C++

1.2.2 Parallel Machine Model

- PRAM
 - a natural extension of the Random Access Machine (RAM) serial architecture
 - consists of p processors and a global memory of unbounded size that is uniformly accessible to all processors
 - processors share a common clock but may execute different instructions in each cycle
- Handling of simultaneous memory accesses
 - Exclusive-read, exclusive-write (EREW)
 - Concurrent-read, exclusive-write (CREW)
 - Exclusive-read, concurrent-write (ERCW)
 - Concurrent-read, concurrent-write (CRCW)

What does concurrent write mean?

Common: write only if all values are identical.

Arbitrary: write the data from a randomly selected processor.

Priority: follow a predetermined priority order.

Sum: write the sum of all data items.

1.3 Different grains of parallelism

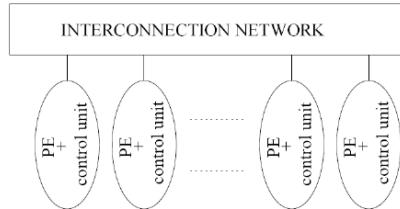

- Granularity: the ratio of computation to communication
 - periods of computation are separated from periods of communication by synchronization events
 - constrained by the inherent characteristics of the used algorithms
 - the parallel programmer must select the right granularity to benefit from the underlying platform
- Chunking
 - determining the amount of data to assign to each task (chunk or grain size)
- Which Granularity will lead to best performance?
 - depends on the algorithm and the used hardware environment
 - general rule: increase grain size if the communication overhead is too large

1.3.1 Trade-offs associated with chunk size

- Fine-grained parallelism
 - low arithmetic intensity
 - may not have enough work to hide long-duration asynchronous communication
 - facilitates load balancing by providing a larger number of more manageable (i.e. smaller) work units
 - too fine granularity can produce slower parallel implementation than the serial execution (too much overhead required for communication)
- Coarse-grained parallelism
 - high arithmetic intensity
 - complete applications can serve as the grain of parallelism
 - more difficult to load balance efficiently

1.4 Control structure of parallel platforms

- SIMD: Single Instruction stream, Multiple Data stream
 - there is a single control unit that dispatches the same instruction to various processors (that work on different data)
 - data parallelization
- MIMD: Multiple Instruction stream, Multiple Data stream
 - each processor has its own control unit
 - each processor can execute different instructions on different data items



1.4.1 SIMD Computers

- Hardware requirements
 - SIMD computers require less HW than MIMD computers (only one control unit)
 - SIMD computers require less memory (only one copy of the program is stored)
- Current implementations
 - Graphics Processing Units (GPUs)
 - Digital Signal Processors (DSPs) are widely used in cameras and sound equipments
 - Co-processing units in Intel CPUs: SSE_x, AVX-512
- Software requirements
 - SIMD relies on the regular structure of computations (such as those in image and video processing or in deep learning)
 - it is often necessary to selectively turn off operations on certain data items


1.4.2 MIMD Computers

- Single Program Multiple Data (SPMD)
 - a simple variant of MIMD executes the same program on different processors
 - SPMD and MIMD are closely related in terms of programming flexibility and underlying architectural support
 - a single program consisting of several programs in a large switch block with conditions specified by the task identifiers is equivalent to the MIMD model
- Current MIMD implementations
 - SPARC servers, multiprocessor PCs, NASA Beowulf inspired workstation clusters
- Key advantages of workstation clusters
 - high performance workstations and PCs available at low cost
 - latest processors can easily be incorporated into the system as they become available
 - existing software can be used or modified

1.5 Communication models of parallel platforms

- Shared-Address-Space Platforms (Multiprocessors)
 - part (or all) of the memory is accessible to all processors
 - processors interact by modifying data objects stored in this shared-address-space
 - uniform or non-uniform memory access time (UMA vs. NUMA)
- Message Passing Platforms (Multicomputers)
 - comprise of a set of processors and their own (exclusive) memory
 - instances come naturally from clustered workstations (distributed systems) and non-shared-address-space multi-computers
 - are programmed using sending messages (variants of send and receive primitives)
 - libraries such as MPI (1990's) provide such primitives

1.6 Interconnection networks



- Interconnection Networks for Parallel Computers
 - carry data between processors and to memory
 - are made of switches and links (wires, fiber)
 - are classified as static or dynamic
 - * static (direct) networks consist of point-to-point communication links among processing nodes
 - * dynamic (indirect) networks are built using switches and communication links
- Network Topologies
 - a variety of network topologies have been proposed and implemented
 - tradeoff performance for cost
 - two basic categories: physical and logical topologies
 - commercial machines often implement hybrids of multiple topologies

1.6.1 Interconnection Network for HPC

- Infiniband
 - a computer-networking communications standard used in HPC that features very high throughput and very low latency
 - it is used for data interconnect both among and within computers
 - it is also utilized as either a direct, or switched interconnect between servers and storage systems, as well as an interconnect between storage systems
 - it is designed to be scalable and uses a switched fabric network topology

Year	FDR 2011	EDR 2014	HDR 2017	NDR 2020	XDR 2023
Throughput, per 1x [Gbit/s]	13.64	25	50	100	250
Speed for 4x links [Gbit/s]	54.54	100	200	400	1000
Speed for 12x links [Gbit/s]	163.64	300	600	1200	3000
Latency [μ s]	0.7	0.5	0.5	tbd	tbd

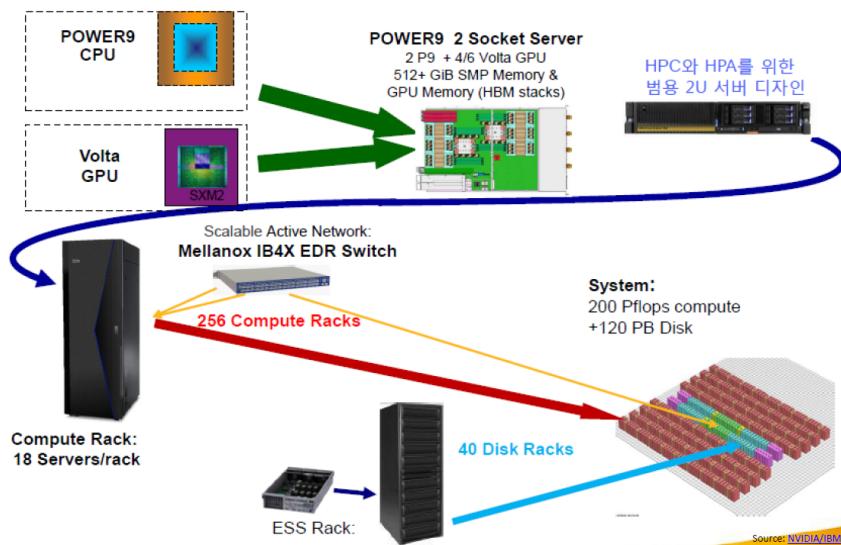
- PCI Express Version 4

- a high-speed serial computer expansion bus standard
- has numerous improvements over the older standards
 - * higher maximum system bus throughput
 - * lower I/O pin count and smaller physical footprint
- has been drafted with final specifications expected in 2017
- throughput:
 - * x1: 1.969 GByte/s
 - * x16: 31.508 GByte/s
- external cabling: Thunderbolt

- NVIDIA NVLink

- is a high-bandwidth, energy-efficient interconnect
- enables ultra-fast communication between the CPU and GPU, and between GPU
- throughput:
 - * version 1 (used in NVIDIA Pascal): x1: 20 GByte/s, x4: 80 GByte/s
 - * version 2 (used in IBM Power9 chip, NVIDIA Volta GPUs): x1: 25 GByte/s, x8: 200 GByte/s

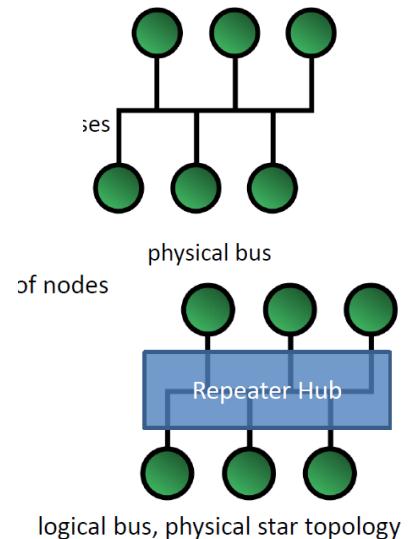
1.6.2 Data-Centric IT Environments



1.7 Network topologies


1.7.1 Network Topologies: Bus

- Principle and Properties
 - some of the simplest and earliest parallel machines used buses
 - all processors access a common bus for exchanging data
 - the distance between any two nodes is $O(1)$ in a bus
- Bottleneck
 - the bandwidth of the shared bus is a major bottleneck
 - typical bus-based machines are limited to dozens of nodes
- Examples
 - WLAN zone (logical bus topology)
 - PCI bus (physical bus topology)




1.7.2 Network Topologies: Star

- Principle and Properties
 - every node is connected only to a common node at the center
 - distance between any pair of nodes is $O(1)$
- Bottleneck
 - the central node
- Example
 - today's Ethernet based LANs with bridging hub (Bridge) or switching hub (Switch) as the center of the star topology



Network Infrastructure: Switching Hub

- Principle and Properties
 - frame forwarding depends on learned physical device-addresses (MAC) per port (Layer-2 switching)
 - non-blocking: several input-output connections can be used in parallel without blocking
 - store-and-forward
 - * the switch buffers and verifies each frame before forwarding it
 - * a frame is received in its entirety before it is forwarded
 - * error checking can be done before forwarding
 - cut-through
 - * the switch starts forwarding after the frame's destination address is received
 - * when the outgoing port is busy at the time, the switch falls back to store-and-forward operation
 - * there is no error checking with this method

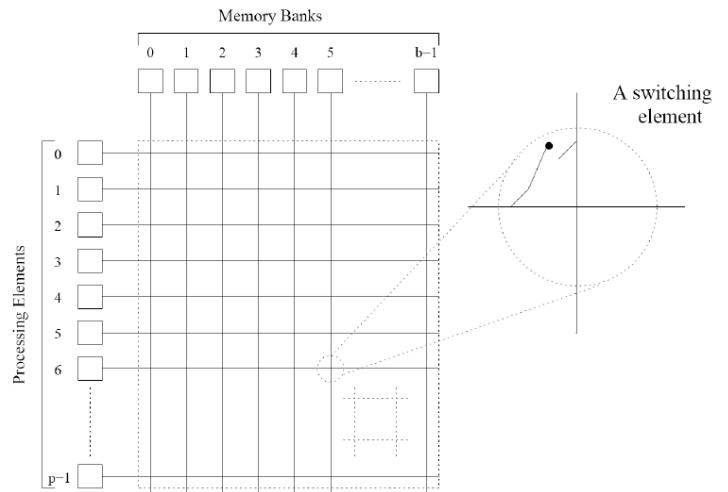


Switching Hub: Advanced Features

- Spanning Tree Protocol → Shortest Path Bridging
 - classic bridges may also interconnect using a spanning tree protocol that disables links so that the resulting local area network is a tree without loops
 - in contrast to routers, spanning tree bridges must have topologies with only one active path between two points
 - IEEE 802.1aq allows all paths to be active with multiple equal cost paths
 - * provides much larger layer 2 topologies (up to 16 million compared to the 4096 VLANs limit)
 - * improves the use of the **mesh topologies** through increased bandwidth and redundancy between all devices by allowing traffic to load share across all paths of a mesh network
- IEEE 802
 - is a family of IEEE standards dealing with local area networks and metropolitan area networks
 - services and protocols specified in IEEE 802 map to the lower two layers (Data Link and Physical) of the seven-layer OSI networking reference model
 - small subset of the working groups
 - * 802.1: higher layer LAN protocols (bridging)
 - * 802.1D: Spanning Tree Protocol (forwarding stopped while the spanning tree re-converged)
 - * 802.1S: Multiple Spanning Tree Protocol
 - * 802.1W: Rapid Spanning Tree Protocol
 - * 802.1aq: Shortest Path Bridging (SPB) (incorporate all the older spanning tree protocols)

1.7.3 Network Topologies: Crossbar


- Principle and Properties
 - a crossbar network uses an $p \cdot b$ grid of switches to connect p inputs to b outputs in a non-blocking manner
- Bottleneck
 - the cost of a crossbar of p processors grows as $O(p^2)$ → difficult to scale for large values of p
- Usage
 - in non-blocking switches
 - between L2- and L2-caches



1.7.4 Network Topologies: Multistage Network


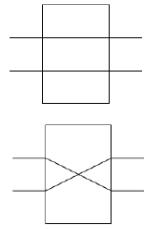
- Scalability
 - busses have excellent cost scalability, but poor performance scalability
 - crossbars have excellent performance scalability but poor cost scalability
 - multistage interconnects strike a compromise between these extremes
- Example: Omega Network
 - it consists of $\log(p)$ stages, where p is the number of inputs/outputs
 - at each stage, input i is connected to output j if:

$$j = \begin{cases} 2i, & 0 \leq i \leq p/2 - 1 \\ 2i + 1 - p, & p/2 \leq i \leq p - 1 \end{cases}$$




Omega Network

- Principle and Properties
 - the perfect shuffle patterns are connected using 2x2 switches
 - the switches operate in two modes: pass-through or cross-over




Linear Array, Mesh, and $k-d$ Mesh

- Principle and Properties
 - in a linear array, each node has two neighbors, one to its left and one to its right
 - if the nodes at either end are connected, we refer to it as a 1-D torus or a ring
 - a generalization to 2 dimensions has nodes with 4 neighbors, to the north, south, east, and west (toroidal mesh)
 - a further generalization to d dimensions has nodes with $2d$ neighbors

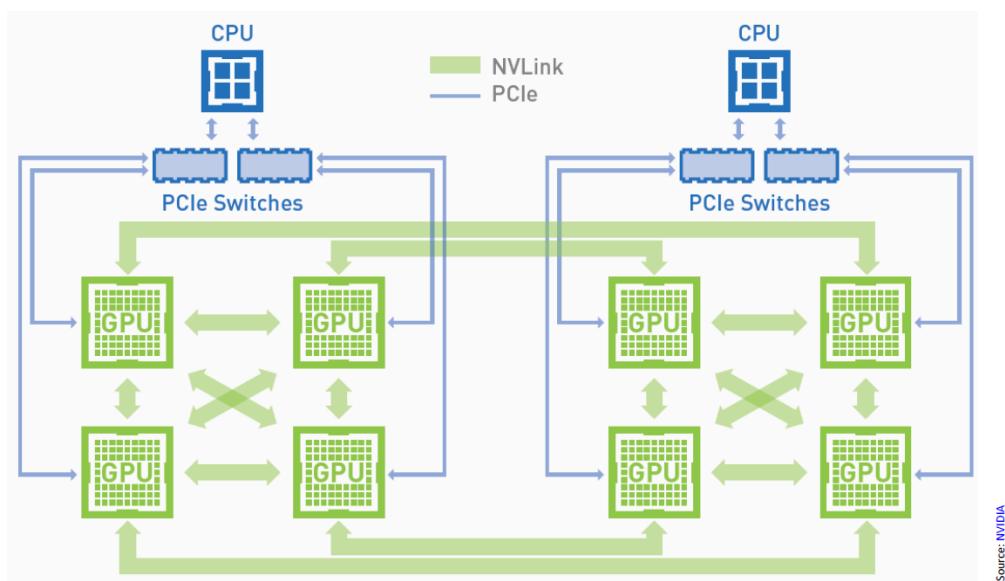


1.7.5 Network Topologies: Hypercube

- Principle and Properties
 - a special case of a d -dimensional mesh is a hypercube
 - $d = \log(p)$, where p is the total number of nodes
 - the distance between any two nodes is at most $\log(p)$
 - each node has $\log(p)$ neighbors
 - the distance between two nodes is given by the number of bit positions at which the two nodes differ




NVIDIA NVLink: Hypercube Mesh Hybrid




Source: NVIDIA

1.7.6 Tree-Based Networks

- Principle and Properties
 - one path between any pair of nodes
 - * linear arrays and star-connected networks are special cases of tree networks
 - the distance between any two nodes is no more than $2\log(p)$
 - links higher up the tree potentially carry more traffic than those at the lower levels
 - trees can be laid out in 2D with no wire crossings
- Fat-Tree



- links higher-up the tree have increased bandwidth



1.7.7 Evaluating Interconnection Networks

- Diameter
 - the distance between the farthest two nodes in the network
- Channel Bandwidth = channel width x channel rate
 - channel width: number of bits that can be communicated simultaneously over a link
 - channel rate: peak data transfer rate per link
- Cross-Section Bandwidth = bisection width x channel bandwidth
 - bisection width: the minimum number of wires one must cut to divide the network into two equal parts
- Cost
 - the number of links or switches (whichever is asymptotically higher) is a meaningful measure of the cost
 - the ability to layout the network
 - the length of wires
- ...

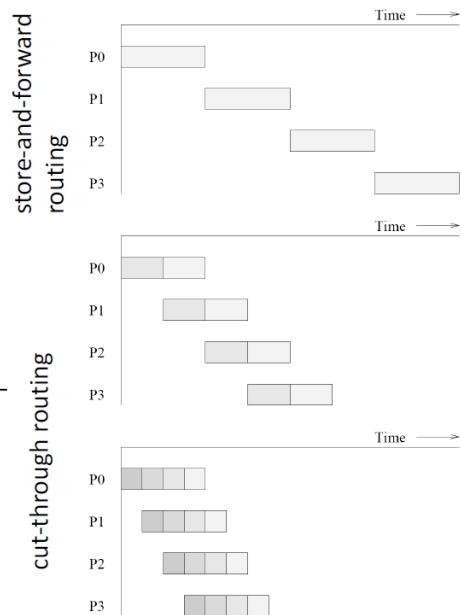
Network	Diameter	Bisection width	Arc connectivity	Cost (No. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$
Star	2	*	1	$p - 1$
Complete binary tree	$2 \cdot \log((p + 1)/2)$	1	1	$p - 1$
Linear array	$p - 1$	1	1	$p - 1$
2D Mesh, no wraparound	$2(\sqrt{p} - 1)$	\sqrt{p}	2	$2(p - \sqrt{p})$
2D wraparound Mesh	$2\lfloor\sqrt{p}/2\rfloor$	$2\sqrt{p}$	4	$2p$
Hypercube	$\log(p)$	$p/2$	$\log(p)$	$(p \cdot \log(p))/2$
Wraparound k -ary d -cube	$d\lfloor k/2 \rfloor$	$2k^{d-1}$	$2d$	dp

* depends on the node (switch) in the center, e.g. Crossbar or Omega Network

Network	Diameter	Bisection width	Arc connectivity	Cost (No. of links)
Crossbar	1	p	1	p^2
Omega Network	$\log(p)$	$p/2$	2	$p/2$
Dynamic Tree	$2 \cdot \log(p)$	1	2	$p - 1$

1.8 Communication costs in parallel systems

- Overhead in parallel programs
 - idling
 - contention
 - communication
- Communication costs depend on
 - communication model
 - the network topology
 - data handling and routing (e.g. packet routing, cut-through routing)
 - associated software protocols
 - ...



1.8.1 Message Passing Costs

- Total time to transfer a message over a network comprises of the following:
 - *Startup time (t_s)*: Time spent at sending and receiving nodes (executing the routing algorithm, programming routers, etc.)
 - *Per-hop time (t_h)*: This time is a function of number of hops and includes factors such as switch latencies, network delays, etc.
 - *Per-word transfer time (t_w)*: This time includes all overheads that are determined by the length of the message. This includes bandwidth of links, error checking and correction, etc.

1.8.2 Cost Model for Communicating Messages

- Communication Costs
 - the cost of communicating a message between two nodes/hops away using cut-through routing is given by

$$t_{\text{comm}} = t_s + l \cdot t_h + m \cdot t_w$$
 - t_h is typically smaller than t_s and t_w , so the second term does not show, when m is large
 - furthermore, it is often not possible to control routing and placement of tasks
- Simplified Cost Model

$$t_{\text{comm}} = t_s + m \cdot t_w$$
- Remarks
 - it is important to note that the original expression for communication time is valid for only uncongested networks
 - if a link takes multiple messages, the corresponding t_w term must be scaled up by the number of messages
 - different communication patterns congest different networks to varying extents

1.8.3 Cost Model for Shared Memory Systems

- Simplified Cost Model (still practical, but accurate cost modeling is more difficult)
 - memory layout is typically determined by the system
 - finite cache sizes can result in cache thrashing

- overheads associated with invalidate and update operations are difficult to quantify
- spatial locality is difficult to model
- pre-fetching can play a role in reducing the overhead associated with data access
- false sharing and contention are difficult to model